

ProblemSet6_Answer

February 20, 2019

0.0.1 Problem Set 6

MACS 30150, Dr. Evans

Due Wednesday, Feb. 20 at 11:30am

Tianxin Zheng

1. Multiple linear regression

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from pandas.plotting import scatter_matrix
import statsmodels.api as sm
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
```

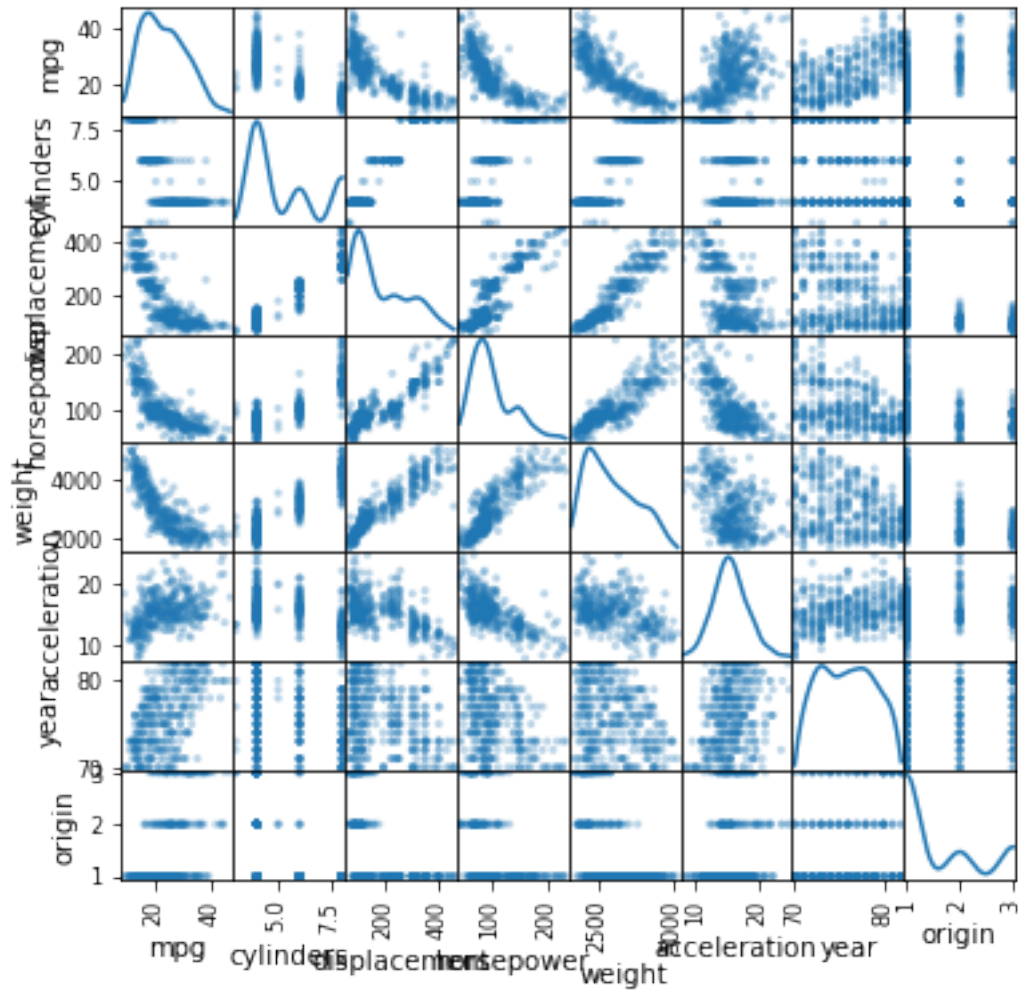
(a)

```
In [2]: df=pd.read_csv("data/Auto.csv", na_values='?')
df.dropna(inplace=True)
```

(b)

```
In [3]: df_quant = df[['mpg', 'cylinders', 'displacement', 'horsepower', 'weight', 'accelerati
```

```
In [4]: scatter_matrix(df_quant, alpha=0.3, figsize=(6, 6), diagonal='kde')
plt.show()
```



(c)

In [5]: `df_quant.corr()`

```
Out [5]:
```

	mpg	cylinders	displacement	horsepower	weight	\
mpg	1.000000	-0.777618	-0.805127	-0.778427	-0.832244	
cylinders	-0.777618	1.000000	0.950823	0.842983	0.897527	
displacement	-0.805127	0.950823	1.000000	0.897257	0.932994	
horsepower	-0.778427	0.842983	0.897257	1.000000	0.864538	
weight	-0.832244	0.897527	0.932994	0.864538	1.000000	
acceleration	0.423329	-0.504683	-0.543800	-0.689196	-0.416839	
year	0.580541	-0.345647	-0.369855	-0.416361	-0.309120	
origin	0.565209	-0.568932	-0.614535	-0.455171	-0.585005	

	acceleration	year	origin
mpg	0.423329	0.580541	0.565209

cylinders	-0.504683	-0.345647	-0.568932
displacement	-0.543800	-0.369855	-0.614535
horsepower	-0.689196	-0.416361	-0.455171
weight	-0.416839	-0.309120	-0.585005
acceleration	1.000000	0.290316	0.212746
year	0.290316	1.000000	0.181528
origin	0.212746	0.181528	1.000000

(d)

```
In [6]: df['const'] = 1
```

```
In [7]: reg1 = sm.OLS(endog=df['mpg'], exog=df[['const', 'cylinders', 'displacement', 'horsepower', 'weight', 'acceleration', 'year', 'origin']],
results1 = reg1.fit()
print(results1.summary())
```

```

OLS Regression Results
=====
Dep. Variable:          mpg      R-squared:                0.821
Model:                  OLS      Adj. R-squared:           0.818
Method:                 Least Squares      F-statistic:           252.4
Date:                   Wed, 20 Feb 2019    Prob (F-statistic):      2.04e-139
Time:                   09:12:02           Log-Likelihood:         -1023.5
No. Observations:       392              AIC:                  2063.
Df Residuals:           384              BIC:                  2095.
Df Model:                7
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	-17.2184	4.644	-3.707	0.000	-26.350	-8.087
cylinders	-0.4934	0.323	-1.526	0.128	-1.129	0.142
displacement	0.0199	0.008	2.647	0.008	0.005	0.035
horsepower	-0.0170	0.014	-1.230	0.220	-0.044	0.010
weight	-0.0065	0.001	-9.929	0.000	-0.008	-0.005
acceleration	0.0806	0.099	0.815	0.415	-0.114	0.275
year	0.7508	0.051	14.729	0.000	0.651	0.851
origin	1.4261	0.278	5.127	0.000	0.879	1.973

```

=====
Omnibus:                 31.906      Durbin-Watson:           1.309
Prob(Omnibus):            0.000      Jarque-Bera (JB):         53.100
Skew:                     0.529      Prob(JB):                 2.95e-12
Kurtosis:                 4.460      Cond. No.                 8.59e+04
=====

```

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 8.59e+04. This might indicate that there are

strong multicollinearity or other numerical problems.

(d.i) The coefficients of 'displacement', 'weight', 'year' and 'origin' are statistically significant at the 1% level.

(d.ii) The coefficients of 'cylinders', 'horsepower' and 'acceleration' are not statistically significant at the 10% level.

(d.iii) Other variables held constant, with 'vehicle year' increasing one unit, 'miles per gallon' will increase around 0.7508 unit.

(e) From the scatterplot matrix from part (b), three variables that look most likely to have a nonlinear relationship with 'mpg' are 'displacement', 'horsepower' and 'weight'.

(e.i)

```
In [8]: df['displacement^2']=np.square(df['displacement'])
        df['horsepower^2']=np.square(df['horsepower'])
        df['weight^2']=np.square(df['weight'])
        df['acceleration^2']=np.square(df['acceleration'])

In [9]: reg2 = sm.OLS(endog=df['mpg'], exog=df[['const', 'cylinders', 'displacement', 'horsepower',
                                                'displacement^2', 'horsepower^2', 'weight^2',
                                                'acceleration^2']],
                      results2 = reg2.fit()
                      print(results2.summary())
```

OLS Regression Results						
=====						
Dep. Variable:	mpg	R-squared:	0.870			
Model:	OLS	Adj. R-squared:	0.866			
Method:	Least Squares	F-statistic:	230.2			
Date:	Wed, 20 Feb 2019	Prob (F-statistic):	1.75e-160			
Time:	09:12:03	Log-Likelihood:	-962.02			
No. Observations:	392	AIC:	1948.			
Df Residuals:	380	BIC:	1996.			
Df Model:	11					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	20.1084	6.696	3.003	0.003	6.943	33.274
cylinders	0.2519	0.326	0.773	0.440	-0.389	0.893
displacement	-0.0169	0.020	-0.828	0.408	-0.057	0.023
horsepower	-0.1635	0.041	-3.971	0.000	-0.244	-0.083
weight	-0.0136	0.003	-5.069	0.000	-0.019	-0.008
acceleration	-2.0884	0.557	-3.752	0.000	-3.183	-0.994

year	0.7810	0.045	17.512	0.000	0.693	0.869
origin	0.6104	0.263	2.320	0.021	0.093	1.128
displacement^2	2.257e-05	3.61e-05	0.626	0.532	-4.83e-05	9.35e-05
horsepower^2	0.0004	0.000	2.943	0.003	0.000	0.001
weight^2	1.514e-06	3.69e-07	4.105	0.000	7.89e-07	2.24e-06
acceleration^2	0.0576	0.016	3.496	0.001	0.025	0.090

Omnibus:	33.614	Durbin-Watson:	1.576
Prob(Omnibus):	0.000	Jarque-Bera (JB):	77.985
Skew:	0.438	Prob(JB):	1.16e-17
Kurtosis:	5.002	Cond. No.	5.13e+08

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 [2] The condition number is large, 5.13e+08. This might indicate that there are strong multicollinearity or other numerical problems.

(e.ii) The adjusted R-squared is 0.866, which is better than 0.818 in the first model.

(e.iii) The coefficients on displacement and its squared term changes from statistically significant at 1% level to nonsignificant at 10% level.

(e.iv) The coefficients on cylinders are not statistically significant at the 10% level.

(f)

```
In [10]: results2.predict(exog=[1, 6, 200, 100, 3100, 15.1, 99, 1, 200**2, 100**2, 3100**2, 15.1**2])
Out[10]: array([38.7321111])
```

The predicted miles per gallon mpg of a car with 6 cylinders, displacement of 200, horsepower of 100, a weight of 3,100, acceleration of 15.1, model year of 1999, and origin of 1, would be 38.73.

2. Classification problem: KNN by hand and in Python

```
In [11]: table=pd.DataFrame({"X1":[0,2,0,0,-1,1], "X2":[3,0,1,1,0,1],
                             "X3":[0,0,3,2,1,1], "Y":["Red", "Red", "Red", "Green", "Green", "Red"]})
table["Eucl.Dist.from X1=X2=X3=0 "]=round(np.sqrt((table["X1"]-0)**2+(table["X2"]-0)**2+(table["X3"]-0)**2),2)
table.index+=1
table
```

```
Out[11]:
```

	X1	X2	X3	Y	Eucl.Dist.from X1=X2=X3=0
1	0	3	0	Red	3.000
2	2	0	0	Red	2.000
3	0	1	3	Red	3.162
4	0	1	2	Green	2.236
5	-1	0	1	Green	1.414
6	1	1	1	Red	1.732

(a) The Euclidean distance between each observation and the test point $X_1 = X_2 = X_3 = 0$ shows below:

$$\begin{aligned}d_1 &= 3 \\d_2 &= 2 \\d_3 &= \sqrt{10} \\d_4 &= \sqrt{5} \\d_5 &= \sqrt{2} \\d_6 &= \sqrt{3}\end{aligned}$$

(b) The KNN prediction with $K = 1$ is Green, because the closest observation to $X_1 = X_2 = X_3 = 0$ is the 5th observation, which is green.

(c) The KNN prediction with $K = 3$ is Red, because the nearest 3 observations to $X_1 = X_2 = X_3 = 0$ are respectively observation 2, 5, 6. Observation 2 and 6 are red and observation 5 is green, so the prediction is red.

(d) If the Bayes (optimal) decision boundary in this problem is highly nonlinear, then we would expect the best value for K to be large. Since the boundary in this problem is highly nonlinear, large K can capture the feature of surrounding points in all directions better. We could better approximate the optimal decision boundary by increasing K .

(e)

```
In [12]: Points = np.array([[0,3,0], [2,0,0], [0,1,3], [0,1,2], [-1,0,1], [1,1,1]])
ys = np.array(['red', 'red', 'red', 'green', 'green', 'red']).reshape(-1, 1)
test_point = np.array([1,1,1]).reshape(1,-1)
knn = KNeighborsClassifier(n_neighbors=2, weights='distance').fit(Points, ys)
print("The KNN classifier of the test point  $X_1 = X_2 = X_3 = 1$  with  $K = 2$  is {}.".format
```

The KNN classifier of the test point $X_1 = X_2 = X_3 = 1$ with $K = 2$ is red.

```
/Users/tianxinzheng/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:4: DataConversionWarning:
  after removing the cwd from sys.path.
```

3. Multivariable logistic (logit) regression

```
In [13]: df['mpg_high'] = np.where(df['mpg'] > np.median(df['mpg']), 1, 0)
```

(a)

```
In [14]: reg3 = sm.Logit(endog=df['mpg_high'], exog=df[['const', 'cylinders', 'displacement'],
results3 = reg3.fit()
print(results3.summary())
```

Optimization terminated successfully.
 Current function value: 0.200944
 Iterations 9

Logit Regression Results

Dep. Variable:	mpg_high	No. Observations:	392
Model:	Logit	Df Residuals:	384
Method:	MLE	Df Model:	7
Date:	Wed, 20 Feb 2019	Pseudo R-squ.:	0.7101
Time:	09:12:03	Log-Likelihood:	-78.770
converged:	True	LL-Null:	-271.71
		LLR p-value:	2.531e-79

	coef	std err	z	P> z	[0.025	0.975]
const	-17.1549	5.764	-2.976	0.003	-28.452	-5.858
cylinders	-0.1626	0.423	-0.384	0.701	-0.992	0.667
displacement	0.0021	0.012	0.174	0.862	-0.021	0.026
horsepower	-0.0410	0.024	-1.718	0.086	-0.088	0.006
weight	-0.0043	0.001	-3.784	0.000	-0.007	-0.002
acceleration	0.0161	0.141	0.114	0.910	-0.261	0.293
year	0.4295	0.075	5.709	0.000	0.282	0.577
origin	0.4773	0.362	1.319	0.187	-0.232	1.187

Possibly complete quasi-separation: A fraction 0.14 of observations can be perfectly predicted. This might indicate that there is complete quasi-separation. In this case some parameters will not be identified.

Coefficients of weight and year are statistically significant at the 5% level.

(b)

```
In [15]: Y = df['mpg_high']
X = df[['cylinders', 'displacement', 'horsepower', 'weight', 'acceleration', 'year',
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size = 0.5, random_state=
```

(c)

```
In [16]: reg4 = LogisticRegression(random_state=10, solver='lbfgs', multi_class='multinomial',
print('The estimated intercept is: ', reg4.intercept_)
print('The estimated coefficient is: ', reg4.coef_[0])
```

The estimated intercept is: [-0.10026869]

The estimated coefficient is: [-0.65773519 0.00857663 -0.01766136 -0.00257161 -0.10958242 0.04697382]

```
In [17]: coef = pd.DataFrame({"coefficient":['constant','cylinders','displacement','horsepower',
                                             "estimate":list(reg4.intercept_) + list(reg4.coef_[0])})
        print(coef)
```

	coefficient	estimate
0	constant	-0.100269
1	cylinders	-0.657735
2	displacement	0.008577
3	horsepower	-0.017661
4	weight	-0.002572
5	acceleration	-0.109582
6	year	0.167356
7	origin	-0.046974

(d)

```
In [18]: y_pred = reg4.predict(X_test)
        cm = confusion_matrix(y_test, y_pred)
        print("Confusion matrix:")
        print(cm)
```

Confusion matrix:

```
[[86 13]
 [12 85]]
```

```
In [19]: print("Classification report:")
        print(classification_report(y_test, y_pred))
```

Classification report:

	precision	recall	f1-score	support
0	0.88	0.87	0.87	99
1	0.87	0.88	0.87	97
avg / total	0.87	0.87	0.87	196

The F1-scores are same. This model predicts equally well on low mpg and high mpg.