

PS7_ANSWER

February 27, 2019

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
#plt.style.use('seaborn')
import statsmodels.api as sm
from sklearn.model_selection import train_test_split
from sklearn import preprocessing
from sklearn.linear_model import LogisticRegression
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split, LeaveOneOut, KFold, cross_val_score
from sklearn import metrics
from sklearn.metrics import classification_report, mean_squared_error
from scipy.interpolate import LSQUnivariateSpline
```

```
In [2]: df = pd.read_csv("data/strongdrink.txt")
df.head()
```

```
Out[2]:
```

	cultivar	alco	malic	ash	alk	magn	tot_phen	flav	nonfl_phen	\
0	1	14.23	1.71	2.43	15.6	127	2.80	3.06	0.28	
1	1	13.20	1.78	2.14	11.2	100	2.65	2.76	0.26	
2	1	13.16	2.36	2.67	18.6	101	2.80	3.24	0.30	
3	1	14.37	1.95	2.50	16.8	113	3.85	3.49	0.24	
4	1	13.24	2.59	2.87	21.0	118	2.80	2.69	0.39	

	proanth	color_int	hue	OD280rat	proline
0	2.29	5.64	1.04	3.92	1065
1	1.28	4.38	1.05	3.40	1050
2	2.81	5.68	1.03	3.17	1185
3	2.18	7.80	0.86	3.45	1480
4	1.82	4.32	1.04	2.93	735

```
In [3]: df["cultivar"].value_counts()
```

```
Out[3]: 2    71
1     59
3     46
Name: cultivar, dtype: int64
```

a. Multinomial logistic regression

```
In [4]: y = df["cultivar"]
        X = df[["alco", "malic", "tot_phen", "color_int"]]
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25, random_state=42)

In [5]: clf = LogisticRegression(solver='newton-cg', multi_class='multinomial').fit(X_train, y_train)

In [6]: pd.DataFrame({"j=1": np.append(clf.intercept_[0], clf.coef_[0]),
                      "j=2": np.append(clf.intercept_[1], clf.coef_[1])},
                      index=["beta0", "beta1", "beta2", "beta3", "beta4"])
```

```
Out[6]:
```

	j=1	j=2
beta0	-24.011332	22.801680
beta1	1.700433	-1.467985
beta2	-0.265610	-0.333051
beta3	1.223894	0.664006
beta4	0.022748	-0.922709

```
In [7]: y_pred = clf.predict(X_test)
        print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
1	0.87	1.00	0.93	13
2	1.00	0.90	0.95	21
3	1.00	1.00	1.00	10
avg / total	0.96	0.95	0.96	44

The error rates are 13%, 0%, and 0% for group 1, 2, and 3, respectively.

The model is best at predicting the third group (highest f1-score).

The one with the most observations is the second group. So the most accurately predicted category is not the one with the most observations.

b. LOOCV

```
In [8]: Xvars = df[['alco', 'malic', 'tot_phen', 'color_int']].values
        yvars = df['cultivar'].values
        N_loo = Xvars.shape[0]
        loo = LeaveOneOut()
        loo.get_n_splits(Xvars)
        MSE_vec = np.zeros(N_loo)
        y_test_lst = np.zeros(N_loo)
        y_pred_lst = np.zeros(N_loo)
```

```
In [9]: for train_index, test_index in loo.split(Xvars):
        X_train, X_test = Xvars[train_index], Xvars[test_index]
        y_train, y_test = yvars[train_index], yvars[test_index]
        LogReg = LogisticRegression()
        LogReg.fit(X_train, y_train)
        y_pred = LogReg.predict(X_test)
        y_pred_lst[test_index] = y_pred
        y_test_lst[test_index] = y_test
        MSE_vec[test_index] = (y_test != y_pred)
        print('MSE for test set', test_index, ' is', MSE_vec[test_index])
```

```
MSE for test set [0] is [0.]
MSE for test set [1] is [0.]
MSE for test set [2] is [0.]
MSE for test set [3] is [0.]
MSE for test set [4] is [0.]
MSE for test set [5] is [0.]
MSE for test set [6] is [0.]
MSE for test set [7] is [0.]
MSE for test set [8] is [0.]
MSE for test set [9] is [0.]
MSE for test set [10] is [0.]
MSE for test set [11] is [1.]
MSE for test set [12] is [0.]
MSE for test set [13] is [0.]
MSE for test set [14] is [0.]
MSE for test set [15] is [0.]
MSE for test set [16] is [0.]
MSE for test set [17] is [0.]
MSE for test set [18] is [0.]
MSE for test set [19] is [0.]
MSE for test set [20] is [0.]
MSE for test set [21] is [1.]
MSE for test set [22] is [1.]
MSE for test set [23] is [1.]
MSE for test set [24] is [1.]
MSE for test set [25] is [1.]
MSE for test set [26] is [0.]
MSE for test set [27] is [1.]
MSE for test set [28] is [0.]
MSE for test set [29] is [0.]
MSE for test set [30] is [0.]
MSE for test set [31] is [0.]
MSE for test set [32] is [1.]
MSE for test set [33] is [0.]
MSE for test set [34] is [1.]
MSE for test set [35] is [0.]
MSE for test set [36] is [0.]
```

MSE for test set [37] is [1.]
MSE for test set [38] is [1.]
MSE for test set [39] is [0.]
MSE for test set [40] is [0.]
MSE for test set [41] is [1.]
MSE for test set [42] is [0.]
MSE for test set [43] is [1.]
MSE for test set [44] is [0.]
MSE for test set [45] is [0.]
MSE for test set [46] is [0.]
MSE for test set [47] is [0.]
MSE for test set [48] is [0.]
MSE for test set [49] is [0.]
MSE for test set [50] is [0.]
MSE for test set [51] is [0.]
MSE for test set [52] is [0.]
MSE for test set [53] is [0.]
MSE for test set [54] is [0.]
MSE for test set [55] is [0.]
MSE for test set [56] is [0.]
MSE for test set [57] is [0.]
MSE for test set [58] is [0.]
MSE for test set [59] is [0.]
MSE for test set [60] is [0.]
MSE for test set [61] is [1.]
MSE for test set [62] is [0.]
MSE for test set [63] is [1.]
MSE for test set [64] is [0.]
MSE for test set [65] is [1.]
MSE for test set [66] is [1.]
MSE for test set [67] is [1.]
MSE for test set [68] is [0.]
MSE for test set [69] is [0.]
MSE for test set [70] is [0.]
MSE for test set [71] is [0.]
MSE for test set [72] is [0.]
MSE for test set [73] is [0.]
MSE for test set [74] is [0.]
MSE for test set [75] is [0.]
MSE for test set [76] is [0.]
MSE for test set [77] is [0.]
MSE for test set [78] is [0.]
MSE for test set [79] is [0.]
MSE for test set [80] is [0.]
MSE for test set [81] is [0.]
MSE for test set [82] is [0.]
MSE for test set [83] is [1.]
MSE for test set [84] is [0.]

MSE for test set [85] is [0.]
MSE for test set [86] is [0.]
MSE for test set [87] is [0.]
MSE for test set [88] is [0.]
MSE for test set [89] is [0.]
MSE for test set [90] is [0.]
MSE for test set [91] is [0.]
MSE for test set [92] is [0.]
MSE for test set [93] is [0.]
MSE for test set [94] is [0.]
MSE for test set [95] is [0.]
MSE for test set [96] is [0.]
MSE for test set [97] is [0.]
MSE for test set [98] is [1.]
MSE for test set [99] is [0.]
MSE for test set [100] is [0.]
MSE for test set [101] is [0.]
MSE for test set [102] is [0.]
MSE for test set [103] is [0.]
MSE for test set [104] is [0.]
MSE for test set [105] is [0.]
MSE for test set [106] is [0.]
MSE for test set [107] is [0.]
MSE for test set [108] is [0.]
MSE for test set [109] is [0.]
MSE for test set [110] is [0.]
MSE for test set [111] is [0.]
MSE for test set [112] is [0.]
MSE for test set [113] is [0.]
MSE for test set [114] is [0.]
MSE for test set [115] is [0.]
MSE for test set [116] is [0.]
MSE for test set [117] is [0.]
MSE for test set [118] is [0.]
MSE for test set [119] is [0.]
MSE for test set [120] is [0.]
MSE for test set [121] is [1.]
MSE for test set [122] is [0.]
MSE for test set [123] is [0.]
MSE for test set [124] is [0.]
MSE for test set [125] is [0.]
MSE for test set [126] is [0.]
MSE for test set [127] is [0.]
MSE for test set [128] is [0.]
MSE for test set [129] is [0.]
MSE for test set [130] is [1.]
MSE for test set [131] is [0.]
MSE for test set [132] is [0.]

```
MSE for test set [133] is [0.]
MSE for test set [134] is [1.]
MSE for test set [135] is [0.]
MSE for test set [136] is [0.]
MSE for test set [137] is [0.]
MSE for test set [138] is [1.]
MSE for test set [139] is [0.]
MSE for test set [140] is [0.]
MSE for test set [141] is [0.]
MSE for test set [142] is [0.]
MSE for test set [143] is [0.]
MSE for test set [144] is [0.]
MSE for test set [145] is [0.]
MSE for test set [146] is [0.]
MSE for test set [147] is [0.]
MSE for test set [148] is [0.]
MSE for test set [149] is [0.]
MSE for test set [150] is [0.]
MSE for test set [151] is [0.]
MSE for test set [152] is [0.]
MSE for test set [153] is [0.]
MSE for test set [154] is [0.]
MSE for test set [155] is [0.]
MSE for test set [156] is [0.]
MSE for test set [157] is [0.]
MSE for test set [158] is [0.]
MSE for test set [159] is [0.]
MSE for test set [160] is [0.]
MSE for test set [161] is [0.]
MSE for test set [162] is [0.]
MSE for test set [163] is [0.]
MSE for test set [164] is [0.]
MSE for test set [165] is [0.]
MSE for test set [166] is [0.]
MSE for test set [167] is [0.]
MSE for test set [168] is [0.]
MSE for test set [169] is [0.]
MSE for test set [170] is [0.]
MSE for test set [171] is [0.]
MSE for test set [172] is [0.]
MSE for test set [173] is [0.]
MSE for test set [174] is [0.]
MSE for test set [175] is [0.]
```

```
In [10]: print(classification_report(y_test_lst, y_pred_lst))

precision    recall  f1-score   support
```

1.0	0.84	0.78	0.81	59
2.0	0.83	0.89	0.86	71
3.0	0.96	0.93	0.95	46
avg / total	0.86	0.86	0.86	176

Error rate for each type is 0.16, 0.17, 0.04 respectively. The most precise estimation is $j=3$.

```
In [11]: MSE_loo = MSE_vec.mean()
        MSE_loo_std = MSE_vec.std()
        print('Test estimate MSE loocv = {}'.format(MSE_loo))
```

Test estimate MSE loocv = 0.13636363636363635.

(c)

```
In [12]: X=df[["alco", "malic", "tot_phen", "color_int"]].values
        y=df["cultivar"].values
        k = 4
        kf = KFold(n_splits=k, random_state=10, shuffle=True)
        kf.get_n_splits(X)

        MSE_vec_kf = np.zeros(k)

        k_ind = int(0)
        for train_index, test_index in kf.split(X):
            print('k index=', k_ind)
            X_train, X_test = X[train_index], X[test_index]
            y_train, y_test = y[train_index], y[test_index]
            LogReg = LogisticRegression(multi_class='multinomial', solver='newton-cg')
            LogReg.fit(X_train, y_train)
            y_pred = LogReg.predict(X_test)
            MSE_vec_kf[k_ind] = (y_test != y_pred).mean()
            print('MSE for test set', k_ind, ' is', MSE_vec_kf[k_ind])
            print(classification_report(y_test, y_pred))
            k_ind += 1
```

k index= 0

MSE for test set 0 is 0.1590909090909091

	precision	recall	f1-score	support
1	0.71	1.00	0.83	12
2	1.00	0.75	0.86	24
3	0.78	0.88	0.82	8
avg / total	0.88	0.84	0.84	44

```

k index= 1
MSE for test set 1 is 0.11363636363636363
      precision    recall  f1-score   support

     1         1.00      0.76      0.87         17
     2         0.69      1.00      0.81         11
     3         1.00      0.94      0.97         16

avg / total         0.92      0.89      0.89         44

```

```

k index= 2
MSE for test set 2 is 0.045454545454545456
      precision    recall  f1-score   support

     1         0.94      1.00      0.97         15
     2         0.94      0.94      0.94         16
     3         1.00      0.92      0.96         13

avg / total         0.96      0.95      0.95         44

```

```

k index= 3
MSE for test set 3 is 0.045454545454545456
      precision    recall  f1-score   support

     1         0.88      1.00      0.94         15
     2         1.00      0.90      0.95         20
     3         1.00      1.00      1.00          9

avg / total         0.96      0.95      0.95         44

```

```

In [13]: print('the average error rate for category 1 is {:.2f}%'.format((1- (0.71 + 1.00 + 0.94)/3)))
          print('the average error rate for category 2 is {:.2f}%'.format((1- (1.00 + 0.69 + 0.94)/3)))
          print('the average error rate for category 3 is {:.2f}%'.format((1- (0.78 + 1.00 + 1.00)/3)))

```

```

the average error rate for category 1 is 11.750000%
the average error rate for category 2 is 9.250000%
the average error rate for category 3 is 5.500000%

```

Compared to those from part (b), the error rates become slightly larger.

Compared to those from part (a), the error rate for category 1 become slightly smaller, but those for category 2 and category 3 become much larger.

```

In [14]: MSE_kf = MSE_vec_kf.mean()
          MSE_kf_std = MSE_vec_kf.std()

```



```
print('test estimate MSE k-fold=', MSE_kf)
print('test estimate MSE standard err=', MSE_kf_std)
```

```
test estimate MSE k-fold= 0.09090909090909091
test estimate MSE standard err= 0.04821182598999188
```

2. Splines and interpolation


(a)

```
In [15]: df1 = pd.read_csv("data/CoolIndex.txt", names=["Age", "Cool"])
df1.head()
```

```
Out[15]:
```

	Age	Cool
0	11.0	10.981602
1	11.0	11.364925
2	11.0	10.190227
3	11.0	9.903725
4	11.0	8.997918

```
In [16]: plt.scatter(x =df1['Age'], y =df1['Cool'], alpha=0.3, label="Points")
plt.title('Scatter Plot of Coolness Index')
plt.xlabel('Age')
plt.ylabel('Coolness Index')
plt.show()
```



(b)

```
In [17]: df1["G1"] = np.where((df1['Age'] >= 11) & (df1['Age'] < 22),1,0)
df1["G2"] = np.where((df1['Age'] >= 22) & (df1['Age'] < 40),1,0)
df1["G3"] = np.where((df1['Age'] >= 40) & (df1['Age'] < 59),1,0)
df1["G4"] = np.where((df1['Age'] >= 59) & (df1['Age'] < 77),1,0)
df1["G5"] = np.where((df1['Age'] >= 77) & (df1['Age'] <= 95),1,0)
```

```
In [18]: X=df1[["G1","G2","G3","G4","G5"]]
reg=sm.OLS(df1['Cool'], X, missing='drop').fit()
print(reg.summary())
```

OLS Regression Results

```
=====
Dep. Variable:          Cool    R-squared:                0.429
Model:                  OLS    Adj. R-squared:             0.427
Method:                 Least Squares    F-statistic:        178.7
Date:                  Wed, 27 Feb 2019    Prob (F-statistic):    3.73e-114
Time:                  10:50:59    Log-Likelihood:       -3214.5
No. Observations:      956    AIC:                  6439.
Df Residuals:          951    BIC:                  6463.
Df Model:              4
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
G1	20.1025	0.562	35.746	0.000	18.999	21.206
G2	34.4758	0.431	80.006	0.000	33.630	35.321
G3	37.6351	0.424	88.814	0.000	36.804	38.467
G4	35.2254	0.485	72.560	0.000	34.273	36.178
G5	27.2964	0.936	29.175	0.000	25.460	29.132

```
=====
Omnibus:                80.102    Durbin-Watson:          1.236
Prob(Omnibus):           0.000    Jarque-Bera (JB):        101.718
Skew:                   0.714    Prob(JB):                8.17e-23
Kurtosis:               3.719    Cond. No.                2.21
=====
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [19]: for i in range(5):
print("Beta_{} = {}".format(i+1, reg.params[i]))
```

Beta_1 = 20.102457252090748


Beta_2 = 34.47578807755938

```
Beta_3 = 37.63510549244961
Beta_4 = 35.22540004024275
Beta_5 = 27.296378244321282
```

```
In [20]: print('The predicted coolness of a 73-year old from the stepwise function is', reg.pa
```

The predicted coolness of a 73-year old from the stepwise function is 35.22540004024275

```
In [21]: plt.scatter(df1['Age'], df1['Cool'], alpha=0.3, label="Points")
plt.plot(df1['Age'], reg.predict(), 'r', label = "OLS Stepwise")
plt.legend()
plt.xlabel("age")
plt.ylabel("Coolness Index")
plt.title("Coolness Index against age")
plt.show()
```



(c)


```
In [22]: df2=df1.groupby('Age').mean()
df2['Age']=df2.index
df2.head()
```

```
Out [22]:
```

	Cool	G1	G2	G3	G4	G5	Age
Age							
11.0	10.110237	1	0	0	0	0	11.0
12.0	9.365623	1	0	0	0	0	12.0
13.0	10.015882	1	0	0	0	0	13.0
14.0	11.747109	1	0	0	0	0	14.0
15.0	15.434739	1	0	0	0	0	15.0

```
In [23]: knots=[22, 40, 59, 77]
linear_spline=LSQUnivariateSpline(df2.Age.values, df2.Cool.values, t=knots, k=1)
```

```
In [24]: plt.scatter(df1.Age, df1.Cool, alpha=0.3, label="Points")
plt.plot(df1.Age, reg.predict(), "r", label = "OLS Stepwise")
plt.plot(df2.Age, linear_spline(df2.Age), "g", label = "Linear Spline")
plt.legend()
plt.xlabel("age")
plt.ylabel("Coolness Index")
plt.title("Coolness Index against age")
plt.show()
```




```
In [25]: print('The predicted coolness of a 73-year old from the linear spline is', linear_spline(73))
```

The predicted coolness of a 73-year old from the linear spline is 32.86784862349653

(d)

```
In [26]: cubic_spline = LSQUnivariateSpline(df2.Age.values, df2.Cool.values, knots, k=3)
```

```
In [27]: plt.scatter(df1.Age, df1.Cool, alpha=0.3, label="Points")
plt.plot(df1.Age, reg.predict(), "r", label = "OLS Stepwise")
plt.plot(df2.Age, linear_spline(df2.Age), "g", label = "Linear Spline")
plt.plot(df2.Age, cubic_spline(df2.Age), "b", label = "Cubic Spline")
plt.legend()
plt.xlabel("age")
plt.ylabel("Coolness Index")
plt.title("Coolness Index against age")
plt.show()
```



```
In [28]: print('The predicted coolness of a 73-year old from the cubic spline is', cubic_spline(73))
```

The predicted coolness of a 73-year old from the cubic spline is 32.642301066279764